




Speedup vs Sequence Length (Float32)



Speedup vs Sequence Length (Float16)



Backward vs Forward Pass Ratio (d_head=64, fp32)



Relative Flash Advantage: FP32 vs FP16

